

Haoxiang Sun

hxiang-sun.github.io · hxiang.sun@gmail.com
GitHub · GScholar · +86 182 555 28272 · +1 650 219 4893
TOEFL: 5.5 [110] (Reading: 5.0, Listening: 6.0, Speaking: 4.5, Writing: 5.5)

Education

2023–2027 Beijing, China	Renmin University of China — B.Eng. Computer Science (Turing Elite Program) + B.Econ. Finance GPA 3.94/4.00 · Rank 1/30 · Research Area Large Language Models Selected Coursework Calculus (96), Adv. Linear Algebra (96), Probability (94), Computer Systems (95)
2026.01 – 2026.08 Berkeley, CA	University of California, Berkeley — Visiting Student · Berkeley NLP / BAIR Advanced NLP (CS 288) — A+, <i>highest score in class (1/150)</i> · Taught by Sewon Min, Alane Suhr Computational Models of Language Use (CS 294-258) — A+ (top 4/25) · Taught by Alane Suhr

Publications

Preprint 2026	Which Models Are Our Models Built On? Auditing Invisible Dependencies in Modern LLMs Sanjay Adhikesaven*, Haoxiang Sun*, Sewon Min* ^{equal contribution} [arXiv · Code · Demo] We present <i>ModSleuth</i> — an agentic system that recursively reconstructs LLM dependency graphs from public artifacts (model & dataset cards, technical reports, code). Recovered 1,060 source-verified dependencies across four major releases, surfacing underreporting, hidden license paths, and train-eval coupling.
ACL 2026 Oral (93 cites)	Challenging the Boundaries of Reasoning: An Olympiad-Level Math Benchmark for LLMs Haoxiang Sun, Yingqian Min, Zhipeng Chen, Wayne Xin Zhao [†] , Ji-Rong Wen [arXiv · Code · Dataset] We present <i>OlymMATH</i> — the first bilingual parallel Olympiad-level benchmark: 200 verifiable problems (English/Chinese) plus 150 Lean formal proofs, sourced from print materials with per-problem expert review. Exposed guessing-based solving in SOTA models (DeepSeek-R1, o3-mini, Gemini 2.5 Pro).
Preprint 2024 (74 cites)	Enhancing LLM Reasoning with Reward-guided Tree Search Jinhao Jiang, ..., Haoxiang Sun, Jia Deng, Wayne Xin Zhao [†] , et al. [arXiv · Code] We present <i>STILL-1</i> — reward-guided tree search; built the Beam-Search baseline and a sympy math verifier.
ACL 2024 Demo (9 cites)	LLMBox: A Comprehensive Library for Large Language Models Tianyi Tang, Yiwen Hu, ..., Haoxiang Sun, ..., Wayne Xin Zhao [†] , Ji-Rong Wen [arXiv · Anthology · Code] We present <i>LLMBox</i> — a unified LLM training/eval platform; built evaluation and few-shot prompting.

Research Experience

2026–Now	NLP Group / BAIR, UC Berkeley — Research Intern · Advisor: Sewon Min Co-first-authored <i>ModSleuth</i> ; built the agentic recursive graph-reconstruction pipeline.
2025–2026	MLL Lab, Northwestern University — Research Intern · Advisor: Manling Li · Mentor: Zihan Wang RAGEN (2.7k+): multi-turn agent context management + Lean env; RL training via veRL. Chain-of-Experts (cyclic MoE): rewrote the MoE-variant pretraining & load-balancing pipeline.
2024–2025	Gaoling School of AI, Renmin University of China — Research Intern · Advisor: Wayne Xin Zhao Led <i>OlymMATH</i> (first author, ACL 2026 Oral); contributed to <i>LLMBox</i> .

Honors & Awards (Selected)

2026	CCF Elite Collegiate Award — 103 awardees nationwide
2026	SenseTime Scholarship — 30 awardees nationwide
2024, 2025	National Scholarship (x2) — Ministry of Education, P.R. China (top ~0.2% nationwide)
2022	National Silver Medal — China Mathematical Olympiad (CMO); ranked 2nd in Anhui Province
2021	National Gold Medal — China Southeast Mathematical Olympiad (CSMO)
2024	Asia Team Champion — Asia Terminal, hosted by Citadel Citadel Securities

Activities (Selected)

2025.12	RUC “Bairong Cup” AI Agent Contest — 1st Place; multimodal interview robot (Technical Lead)
2025.03	ETH Hangzhou Hackathon — MCP multi-agent on-chain data analysis & trading [Demo]
2025.03	Discover Citadel Citadel Securities — visited HK office; Team Champion, Cryptography Challenge
2024–2025	Beijing College Student Innovation Program — Technical Lead, News-Image Generation Teacher-student pipeline: Qwen2.5-72B data synthesis, then LoRA + PPO on a 7B student model.